


PREDICT



@CHELSEA PARLETT




# Based Models

Dr. Chelsea Parlett-Pelleriti

# Decision Trees


I BE-LEAF




IN YOU

@CHELSEA PARLETT


# Tree Vocabulary




# Tree Vocabulary




# Tree Vocabulary



# Tree Vocabulary




# Tree Vocabulary






# Twenty Questions

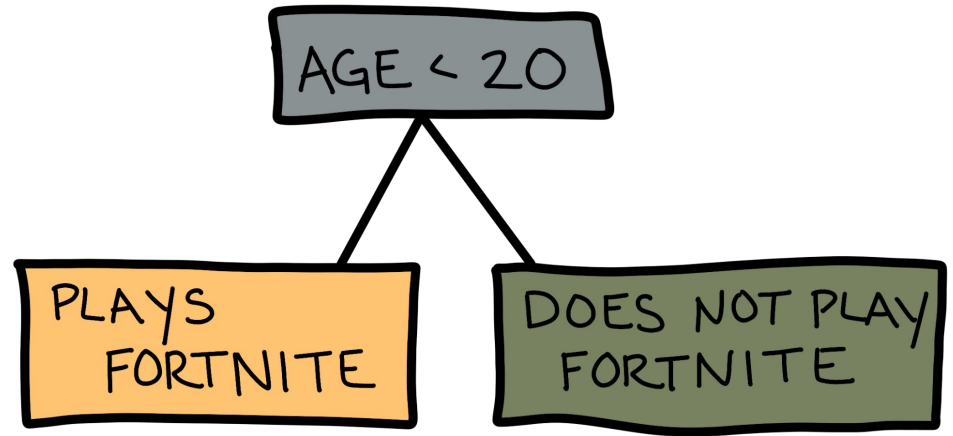

# Simple Tree



## More Complicated Tree



# Data Types




# Gini Impurity and Entropy

$$\text{GI} = 1 - \sum_{i=1}^n (p_i)^2$$

$$\text{E} = - \sum_{i=1}^n p_i * \log(p_i)$$

Goal: choose split so that GI (or entropy) is minimized

# Gini Impurity and Entropy



$$GI = 1 - \sum_{i=1}^n (p_i)^2$$


$$E = - \sum_{i=1}^n p_i * \log(p_i)$$

Figure from the Elements of Statistical Learning (Hastie et al)

# Categorical

$$1 - \sum_{i=1}^n (p_i)^2$$


cats	pet	wfh	children	income
1	0	1	1	34
0	1	0	1	58.3
1	1	1	0	71.5
0	0	0	1	74.9
0	0	0	1	75.3
1	0	0	1	75.6
0	0	0	1	81
1	1	1	0	82.3
1	1	1	0	85.6
1	1	1	1	95.4



# Categorical

$$1 - \sum_{i=1}^n (p_i)^2$$

cats	pet	wfh	children	income
1	0	1	1	34
0	1	0	1	58.3
1	1	1	0	71.5
0	0	0	1	74.9
0	0	0	1	75.3
1	0	0	1	75.6
0	0	0	1	81
1	1	1	0	82.3
1	1	1	0	85.6
1	1	1	1	95.4






# Continuous

$$1 - \sum_{i=1}^n (p_i)^2$$


cats	pet	wfh	children	income
1	0	1	1	34
0	1	0	1	58.3
1	1	1	0	71.5
0	0	0	1	74.9
0	0	0	1	75.3
1	0	0	1	75.6
0	0	0	1	81
1	1	1	0	82.3
1	1	1	0	85.6
1	1	1	1	95.4



# Continuous

$$1 - \sum_{i=1}^n (p_i)^2$$

cats	pet	wfh	children	income
1	0	1	1	34
0	1	0	1	58.3
1	1	1	0	71.5
0	0	0	1	74.9
0	0	0	1	75.3
1	0	0	1	75.6
0	0	0	1	81
1	1	1	0	82.3
1	1	1	0	85.6
1	1	1	1	95.4




# Basic Steps

1. Calculate Gini Impurity (or Entropy/Information Gain) for each node
2. Choose Node with lowest score
3. If the parent node has the lowest score, it is a leaf.

# Variable Importance

1. How much does this feature reduce node impurity?
2. If we shuffle the values of this feature, how much does it reduce the performance?

$$\underbrace{\text{imp}_j}_{\text{importance of node } j} = \underbrace{w_j C_j}_{\text{weighted parent node impurity}} - \underbrace{(w_{\text{left}_j} C_{\text{left}_j} + w_{\text{right}_j} C_{\text{right}_j})}$$




$$fi_i = \frac{\sum_{j \in S_i} \text{imp}_j}{\sum_{k \in S_{all}} \text{imp}_k}; S_i \text{ is set of all nodes that split on feature}_i$$

# Variable Importance

1. How much does this feature reduce node impurity?
2. If we shuffle the values of this feature, how much does it reduce the performance?

$$\underbrace{\text{imp}_j}_{\text{importance of node } j} = \underbrace{w_j C_j}_{\text{weighted parent node impurity}} - \underbrace{(w_{\text{left}_j} C_{\text{left}_j} + w_{\text{right}_j} C_{\text{right}_j})}$$



$$fi_i = \frac{\sum_{j \in S_i} \text{imp}_j}{\sum_{k \in S_{all}} \text{imp}_k}; S_i \text{ is set of all nodes that split on feature}_i$$

# Variable Importance

1. How much does this feature reduce node impurity?
2. If we shuffle the values of this feature, how much does it reduce the performance?

X	Y

# Variable Importance

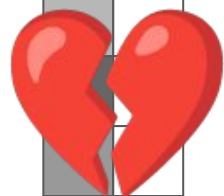
1. How much does this feature reduce node impurity?
2. If we shuffle the values of this feature, how much does it reduce the performance?

X	Y

# Variable Importance

1. How much does this feature reduce node impurity?
2. If we shuffle the values of this feature, how much does it reduce the performance?


X	Y






# Regression Trees

# Regression Trees



# Regression Trees

Instead of checking  
if splits decrease  
Gini Impurity, we  
check if they  
decrease MSE



# Random Forests



**ONE**

**vs.**




**MANY**

# Random Forests

- Bootstrap Aggregating (**Bagging**)
- **Random Feature Selection**

# Random Forests

## BOOTSTRAPPING



@CHELSEA PARLETT

# Random Forests

- Bootstrap Aggregating (**Bagging**)
- **Random Feature Selection**



# Random Forests



# Random Forests



# Random Forests

## Important Hyperparameters


- # of trees
- # of features per tree

# Gradient Boosting Trees


# Gradient Boosting Trees



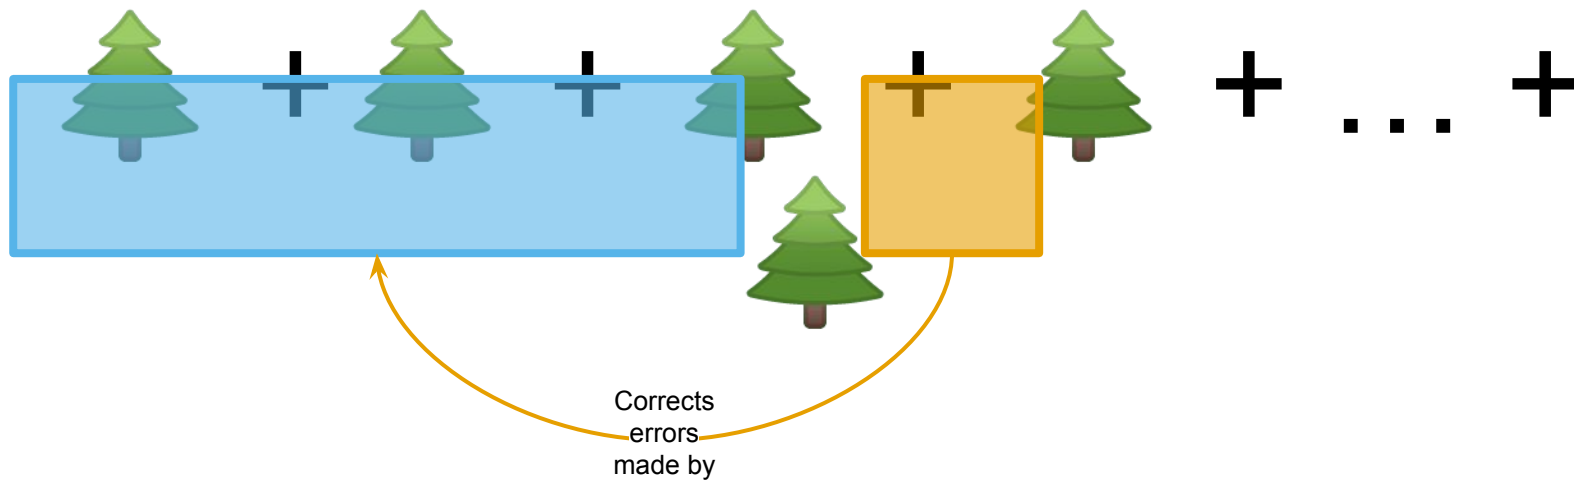
# Gradient Boosting Trees



# Gradient Boosting Trees




# Gradient Boosting Trees





# Gradient Boosting Trees



# Gradient Boosting Tree

	Age	Initial Guess	Residual
Person 1	20		
Person 2	19		
Person 3	21		
Person 4	20		

# Gradient Boosting Tree

	Age	Initial Guess	Residual
Person 1	20	20	
Person 2	19	20	
Person 3	21	20	
Person 4	20	20	

# Gradient Boosting Tree


	Age	Initial Guess	Residual
Person 1	20	20	0
Person 2	19	20	-1
Person 3	21	20	1
Person 4	20	20	0

# Gradient Boosting Tree


$$\text{Actual Value} = \text{Prediction} + \text{Residual}$$

	Age	Initial Guess	Residual
Person 1	20	20	0
Person 2	19	20	-1
Person 3	21	20	1
Person 4	20	20	0


# Gradient Boosting Trees



# Gradient Boosting Trees



# Gradient Boosting Trees





# Gradient Boosting Trees



# Gradient Boosting Trees: Some Extra Math

$$z_i = -\frac{\partial L(y, F_i)}{\partial F_i}$$

Let's generalize that math!  
Having future trees predict  
the error works for a  
regression tree using  $(y-p)^2$   
as it's loss.

# Gradient Boosting Trees: Some Extra Math

$$z_i = -\frac{\partial L(y, F_i)}{\partial F_i}$$

But in all cases, subsequent trees will predict the **negative gradient** ( $z_i$ ) of the Loss with respect to the Ensemble prediction

# Gradient Boosting Trees: Some Extra Math

$$z_i = -\frac{\partial L(y, F_i)}{\partial F_i}$$


But in all cases, subsequent trees will predict the **negative gradient** ( $z_i$ ) of the Loss with respect to the Ensemble prediction

$$F_i = \sum_{t=0}^i z_t$$

$$F_2 = z_0 + z_1 + z_2$$

Baseline prediction

Errors



# Gradient Boosting Trees: Some Extra Math

$$z_i = -\frac{\partial L(y, F_i)}{\partial F_i}$$

But in all cases, subsequent trees will predict the **negative gradient** ( $z_i$ ) of the Loss with respect to the Ensemble prediction

$$F_i = \sum_{t=0}^i z_t$$

$$F_1 = z_0 + z_1$$

$$F_2 = z_0 + z_1 + z_2$$

$$F_i = F_{i-1} + z_i$$

# Gradient Boosting Trees: Some Extra Math

$$z_i = - \frac{\partial L(y, F_i)}{\partial F_i}$$

Negative Gradient of Loss w.r.t.  
Ensemble Prediction

# Gradient Boosting Trees: Some Extra Math

$$z_i = -\frac{\partial L(y, F_i)}{\partial F_i}$$

Negative Gradient of Loss w.r.t.  
Ensemble Prediction

The Negative Gradient tell us what *adjustments* we should make to our prediction ( $F_i$ ) in order to *decrease* our loss.

# Gradient Boosting Trees: Some Extra Math

$$z_i = \frac{\partial L(y, F_i)}{\partial F_i}$$

Negative Gradient of Loss w.r.t.  
Ensemble Prediction

The Negative Gradient tell us what *adjustments* we should make to our prediction ( $F_i$ ) in order to *decrease* our loss.

$$L(y, p) = (y - p)^2$$

Let's choose Squared Error as our Loss

$$-\frac{\partial L(y, p)}{\partial p} = 2(y - p)$$

Negative Gradient



# Gradient Boosting Trees: Some Extra Math

$$z_i = -\frac{\partial L(y, F_i)}{\partial F_i}$$

Negative Gradient of Loss w.r.t.  
Ensemble Prediction

The Negative Gradient tell us what *adjustments* we should make to our prediction ( $F_i$ ) in order to *decrease* our loss.

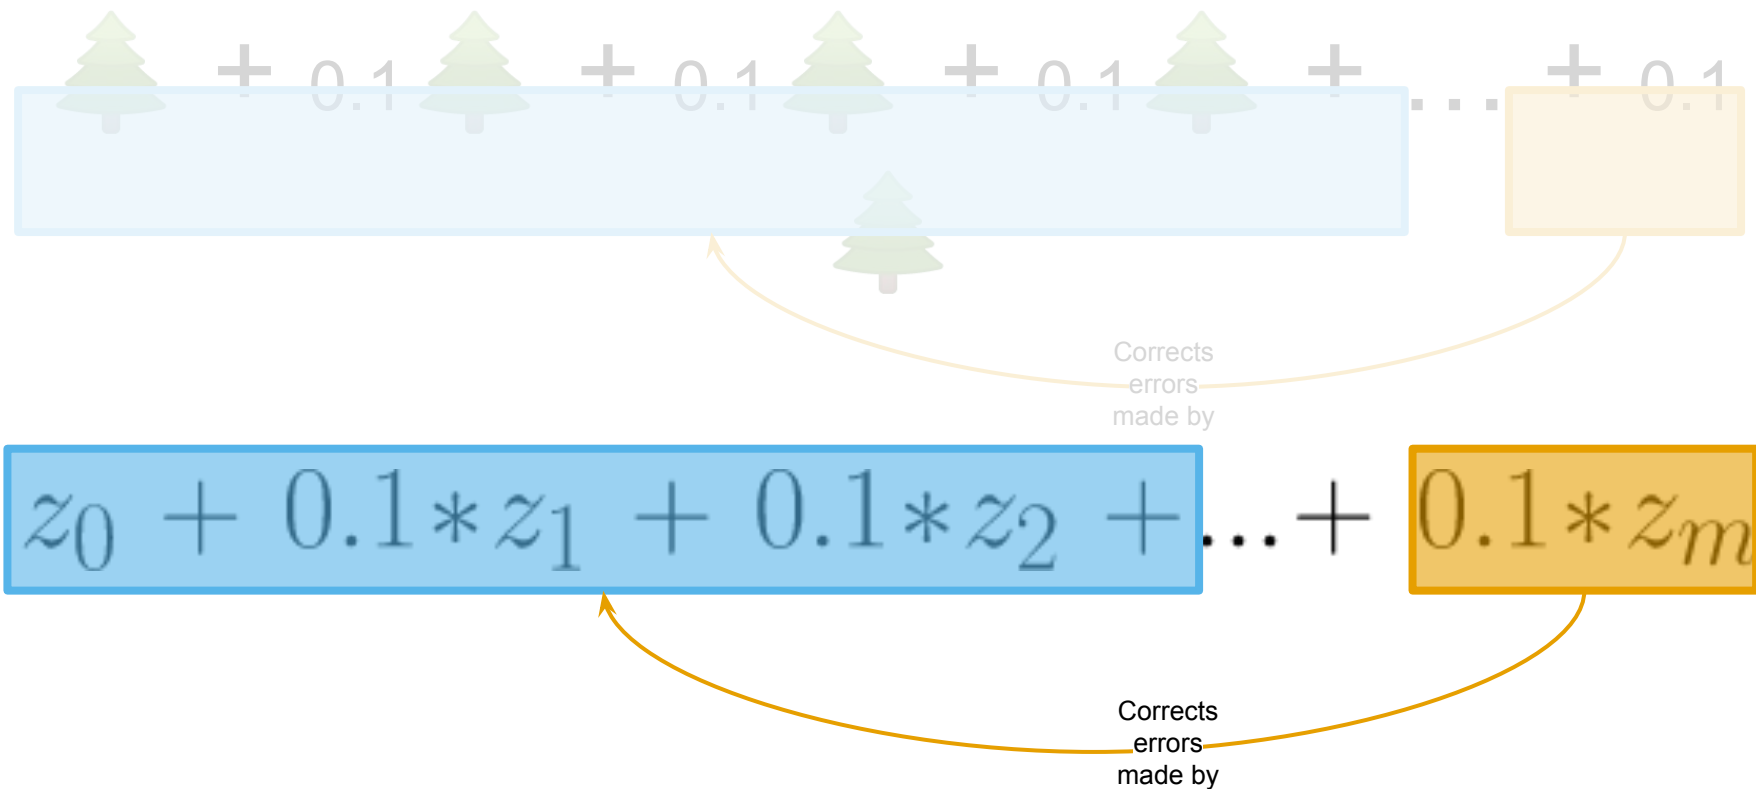
$$L(y, p) = (y - p)^2$$

Let's choose Squared Error as our Loss

$$-\frac{\partial L(y, p)}{\partial p} = 2(y - p)$$

Just the Error!

# Gradient Boosting Trees



# Gradient Boosting Trees: Some Extra Math

$$z_i = -\frac{\partial L(y, F_i)}{\partial F_i}$$

With squared loss, **error is the negative gradient**, but the negative gradient will work in other situations