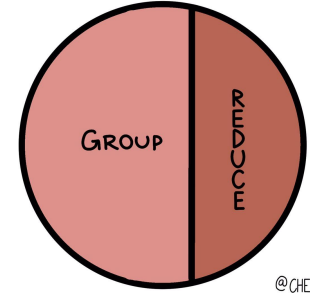


SIMPLIFY




@CHELSEAPARLETT


Hierarchical (Agglomerative) Clustering

Dr. Chelsea Parlett-Pelleriti


HAC




HAC




HAC




HAC




HAC




HAC




Algorithm




Algorithm




Algorithm



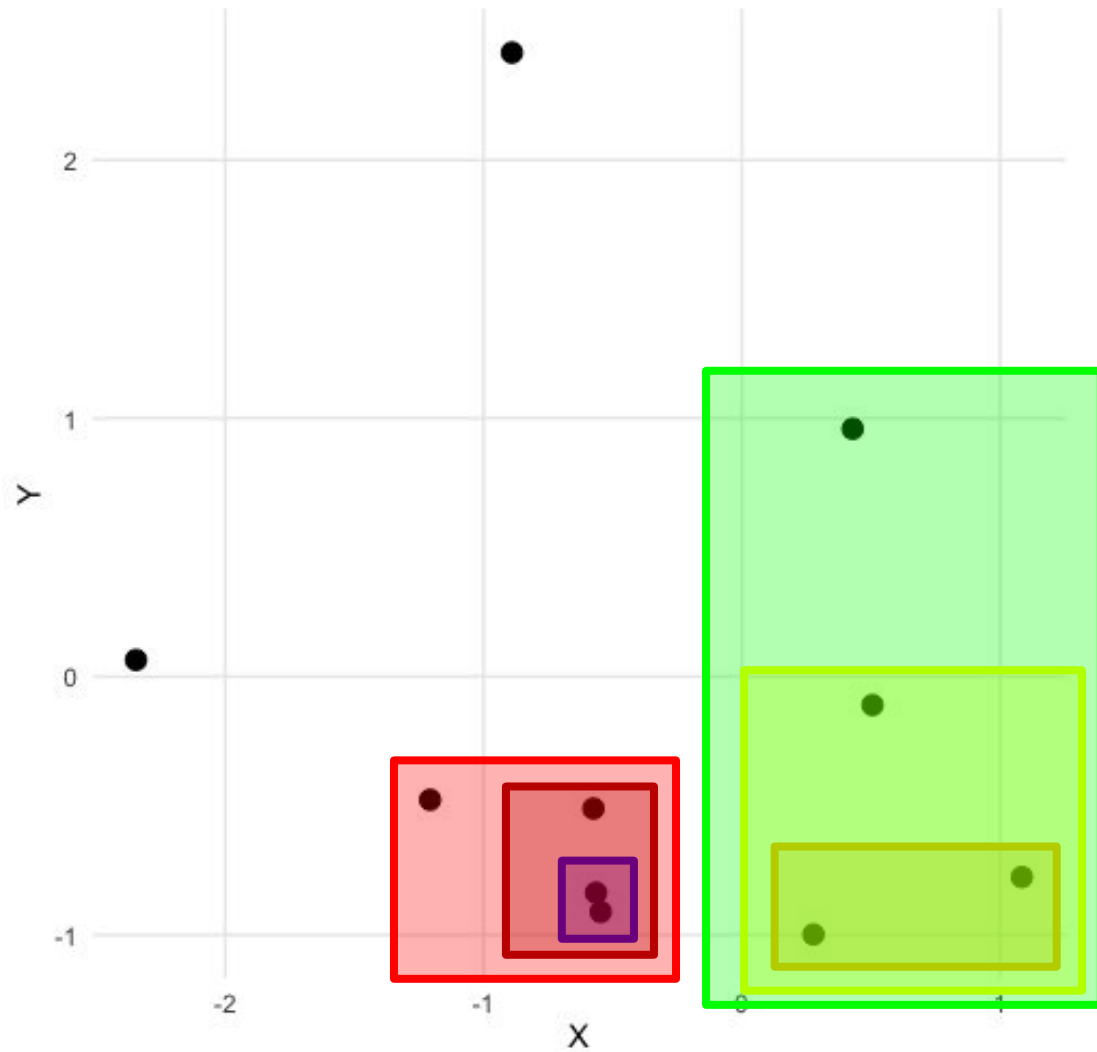
Algorithm




Algorithm




Algorithm




Algorithm



Algorithm



Algorithm




Distance Metrics


- **Euclidean:** Continuous Data
- **Manhattan:** High Dimensions
- **Hamming:** Categories
- **Cosine:** Word Counts

Linkage Criteria

Single




Linkage Criteria




Complete

Linkage Criteria




Average

Linkage Criteria



Linkage Criteria




Ward's


A

B

Reading a Dendrogram



Reading a Dendrogram



HAC


Pros

- Flexible # of clusters
- Model relationship between clusters (hierarchy)
- Flexibility with linkage

Cons

- Very Slow $O(n^3)$
- Cannot un-merge clusters

My Master's Thesis



Application




Fig 5: Visualization of Classification.

Application



Fig. 3: Grouping results. The samples in the 1st & 2nd lines are from HMDB51 and UCF101 respectively.

Application




Fig. 5. Clusters computed by HAC on 55k Data Points

Complexity [\[edit \]](#)

The standard algorithm for **hierarchical agglomerative clustering** (HAC) has a [time complexity](#) of $\mathcal{O}(n^3)$ and requires $\Omega(n^2)$ memory, which makes it too slow for even medium data sets. However, for some special cases, optimal efficient agglomerative methods (of complexity $\mathcal{O}(n^2)$) are known: **SLINK**^[2] for [single-linkage](#) and CLINK^[3] for [complete-linkage clustering](#). With a [heap](#), the runtime of the general case can be reduced to $\mathcal{O}(n^2 \log n)$, an improvement on the aforementioned bound of $\mathcal{O}(n^3)$, at the cost of further increasing the memory requirements. In many cases, the memory overheads of this approach are too large to make it practically usable.

Divisive clustering with an exhaustive search is $\mathcal{O}(2^n)$, but it is common to use faster heuristics to choose splits, such as [k-means](#).